

Programming Assignment 2

Part A: Pooling and Upsampling

A.1 Implementation of PoolUpsampleNet

```
[7]
class PoolUpsampleNet(nn.Module):
    def __init__(self, kernel, num_filters, num_colours, num_in_channels):
        super().__init__()

        # Useful parameters
        padding = kernel // 2

        ##### YOUR CODE GOES HERE #####
        self.firstLayer = nn.Sequential(
            nn.Conv2d(num_in_channels, num_filters, kernel_size=kernel, padding=padding),
            nn.MaxPool2d(2),
            nn.BatchNorm2d(num_filters),
            nn.ReLU()
        )

        self.secondLayer = nn.Sequential(
            nn.Conv2d(num_filters, 2*num_filters, kernel_size=kernel, padding=padding),
            nn.MaxPool2d(2),
            nn.BatchNorm2d(2*num_filters),
            nn.ReLU()
        )



        self.thirdLayer = nn.Sequential(
            nn.Conv2d(2*num_filters, num_filters, kernel_size=kernel, padding=padding),
            nn.Upsample(scale_factor=2),
            nn.BatchNorm2d(num_filters),
            nn.ReLU()
        )

        self.fourthLayer = nn.Sequential(
            nn.Conv2d(num_filters, num_colours, kernel_size=kernel, padding=padding),
            nn.Upsample(scale_factor=2),
            nn.BatchNorm2d(num_colours),
            nn.ReLU()
        )

        self.lastLayer = nn.Conv2d(num_colours, num_colours, kernel, padding=padding)
        #####

    def forward(self, x):
        ##### YOUR CODE GOES HERE #####
        first = self.firstLayer(x)
        second = self.secondLayer(first)
        third = self.thirdLayer(second)
        fourth = self.fourthLayer(third)
        return self.lastLayer(fourth)
        #####
```

A.2 Training Result



The shown figure is the result obtained from the main training loop of PoolUpsampleNet. It does not look good to me as the images are quite blurry with greyscale pixels present. Also, after 25 epoch, accuracy is around 41.4%, which is quite low.

A.3

Assume the kernel size is k

- when each input dimension (width/height) is not doubled (original input)
 - number of weights =
 - number of outputs =
 - number f connections =
- when each input dimension (width/height) is doubled
 - number of weights =
 - number of outputs =
 - number f connections =

Part B: Strided and Transposed Dilated Convolutions

B.1 Implementation of ConvTransposeNet

```

class ConvTransposeNet(nn.Module):
    def __init__(self, kernel, num_filters, num_colours, num_in_channels):
        super().__init__()

        # Useful parameters
        stride = 2
        padding = kernel // 2
        output_padding = 1

        ##### YOUR CODE GOES HERE #####
        self.firstLayer = nn.Sequential(
            nn.Conv2d(num_in_channels, num_filters, kernel_size=kernel, padding=padding, stride=stride),
            nn.BatchNorm2d(num_filters),
            nn.ReLU()
        )

        self.secondLayer = nn.Sequential(
            nn.Conv2d(num_filters, 2*num_filters, kernel_size=kernel, padding=padding, stride=stride),
            nn.BatchNorm2d(2*num_filters),
            nn.ReLU()
        )

        self.thirdLayer = nn.Sequential(
            nn.ConvTranspose2d(2*num_filters, num_filters, kernel_size=kernel, padding=padding, stride=stride,
                               output_padding=output_padding),
            nn.BatchNorm2d(num_filters),
            nn.ReLU()
        )


        self.fourthLayer = nn.Sequential(
            nn.ConvTranspose2d(num_filters, num_colours, kernel_size=kernel, padding=padding, stride=stride,
                               output_padding=output_padding),
            nn.BatchNorm2d(num_colours),
            nn.ReLU()
        )


        self.lastLayer = nn.Conv2d(num_colours, num_colours, kernel_size=kernel, padding=padding)
        #####

    def forward(self, x):
        ##### YOUR CODE GOES HERE #####
        first = self.firstLayer(x)
        second = self.secondLayer(first)
        third = self.thirdLayer(second)
        fourth = self.fourthLayer(third)
        return self.lastLayer(fourth)
        #####

```

B.2





B.3

The trained result seems better from what we got from Part A because 1) the image seems less blurry and 2) the validation accuracy increases from 41.4% to 54.7%.

B.4

B.5

Part C: Skip Connections

C.1 Implementation of UNet

```

class UNet(nn.Module):
    def __init__(self, kernel, num_filters, num_colours, num_in_channels):
        super().__init__()

        # Useful parameters
        stride = 2
        padding = kernel // 2
        output_padding = 1

        ##### YOUR CODE GOES HERE #####
        self.firstLayer = nn.Sequential(
            nn.Conv2d(num_in_channels, num_filters, kernel_size=kernel, padding=padding, stride=stride),
            nn.BatchNorm2d(num_filters),
            nn.ReLU()
        )

        self.secondLayer = nn.Sequential(
            nn.Conv2d(num_filters, 2*num_filters, kernel_size=kernel, padding=padding, stride=stride),
            nn.BatchNorm2d(2*num_filters),
            nn.ReLU()
        )

        self.thirdLayer = nn.Sequential(
            nn.ConvTranspose2d(2*num_filters, num_filters, kernel_size=kernel, padding=padding, stride=stride,
                               output_padding=output_padding),
            nn.BatchNorm2d(num_filters),
            nn.ReLU()
        )

        self.fourthLayer = nn.Sequential(
            nn.ConvTranspose2d(2*num_filters, num_colours, kernel_size=kernel, padding=padding, stride=stride,
                               output_padding=output_padding),
            nn.BatchNorm2d(num_colours),
            nn.ReLU()
        )


        self.lastLayer = nn.Conv2d(num_in_channels + num_colours, num_colours, kernel_size=kernel,
                                     padding=padding)


        #####

    def forward(self, x):
        ##### YOUR CODE GOES HERE #####
        first = self.firstLayer(x)
        second = self.secondLayer(first)
        third = self.thirdLayer(second)
        fourth = self.fourthLayer(torch.cat([third, first], 1))
        return self.lastLayer(torch.cat([fourth, x], 1))
        #####

```

C.2





C.3

Part D: Object Detection

D.1 Fine-tuning from pre-trained models for object detection


```
# Freeze
freeze = [f'model.{x}.' for x in range(freeze)] # layers to freeze
for k, v in model.named_parameters():
    # --- YOUR CODE GOES HERE ---
    if any(x in k for x in freeze):
        v.requires_grad = False
    else:
        v.requires_grad = True
    # -----
```


D.2 Implement the classification loss

D.2.1

```
95 # Define loss criteria
96 # --- YOUR CODE GOES HERE ---
97 BCEcls = nn.BCEWithLogitsLoss(pos_weight=torch.tensor([h['cls_pw']], device=device))
98 # -----
99 BCEobj = nn.BCEWithLogitsLoss(pos_weight=torch.tensor([h['obj_pw']], device=device))
100
101 # Class label smoothing https://arxiv.org/pdf/1902.04103.pdf eqn 3
102 self.cp, self.cn = smooth_BCE(eps=h.get('label_smoothing', 0.0)) # positive, negative BCE targets
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145 # Classification
146 if self.nc > 1: # cls loss (only if multiple classes)
147     t = torch.full_like(ps[:, 5:], self.cn, device=device) # targets
148     t[range(n), tccls[i]] = self.cp
149     # --- YOUR CODE GOES HERE ---
150     lcls += self.BCEcls(ps[:, 5:], t)
151     # -----
152
```


D.2.2





bus 0.94

bicycle 0

person 0.93

person 0.90

person 0

tie 0.50

person 0.55